# RunPod Proof of Use

CareerBridge AI (careerbridge_ai) | Team CodeCatalyst

A genuine request was submitted to a RunPod Serverless endpoint and completed successfully. The endpoint status returned to Ready with zero running workers after the request.

| Record | Observed value |
|---|---|
| Provider | RunPod Serverless |
| Endpoint | 446e2rl998rklk |
| Endpoint page | RunPod console, Requests tab |
| Deployment | vLLM v2.29.0, model Qwen/Qwen3-0.6B |
| GPU | NVIDIA L4, one worker |
| Request ID | 9488922a-3873-4c59-b42d-8fce867e3bda-e1 |
| Result | Completed; response included model output |
| Prompt / completion tokens | 11 / 16 (27 total) |
| Queue delay / execution time | 153,857 ms / 491 ms |
| Run date | 2026-10-24 |
| Post-request status | Ready; 0 running workers; $0.00000/s |

### Scope of this evidence

This verifies a real inference request through the RunPod console and records the response metadata. It does not claim that the public web application was deployed or that its server-side adapter was exercised end to end. The application adapter is included in the repository and requires private server-side credentials and a database configuration.

The request used a short test prompt. No API key or other secret is included in this document. RunPod usage and billing records in the provider console remain authoritative; this record does not state an exact charge.